

Understanding and Overcoming Parallelism Bottlenecks in ForkJoin Applications



[@gustavopinto](https://twitter.com/gustavopinto)



A. Canino



F. Castor



G. Xu



Y. D. Liu



BINGHAMTON
UNIVERSITY
STATE UNIVERSITY OF NEW YORK




BINGHAMTON
UNIVERSITY
STATE UNIVERSITY OF NEW YORK

Modern Java applications run on parallel architectures



java.lang.Thread

- **Widely used**
- **Low level API**
- **Error prone**



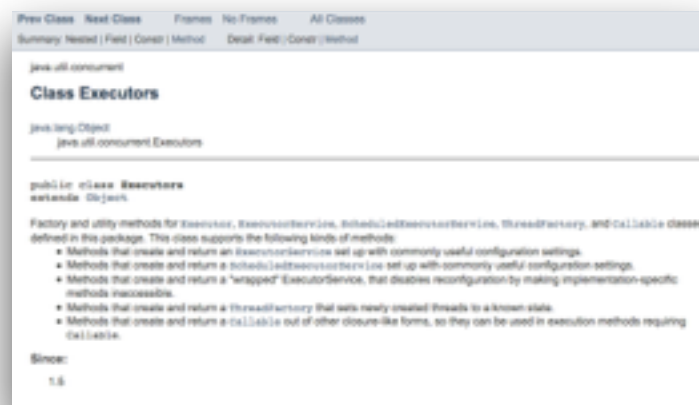
java.util.concurrent.Executors

- **Well used**
- **High Level API**
- **User friendly**

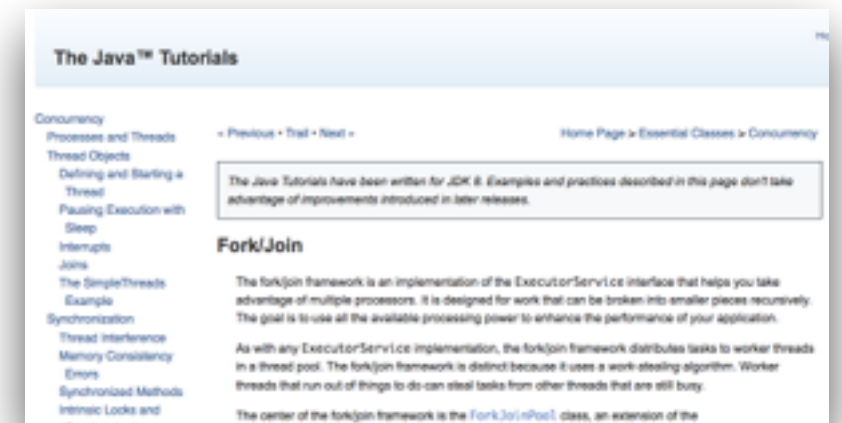
Modern Java applications run on parallel architectures



`java.lang.Thread`



`java.util.concurrent.Executors`



ForkJoin

- Widely used
- Low level API
- Error prone

- Well used
- High Level API
- User friendly

- Can be more used
- Sophisticated API
- Sophisticated scheduler

Modern Java applications run on parallel architectures



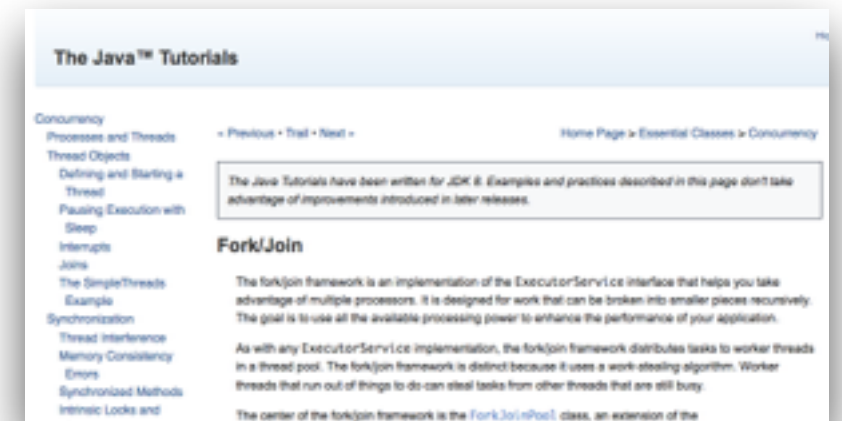
java.lang.Thread

- Widely used
- Low level API
- Error prone



java.util.concurrent.Executors

- Well used
- High Level API
- User friendly




ForkJoin

- Can be more used
- Sophisticated API
- Sophisticated scheduler


Divider and conquer algorithm

Why
ForkJoin?




Divider and conquer algorithm

Why
ForkJoin?




Divider and conquer algorithm

Why
ForkJoin?




Divider and conquer algorithm

Why
ForkJoin?



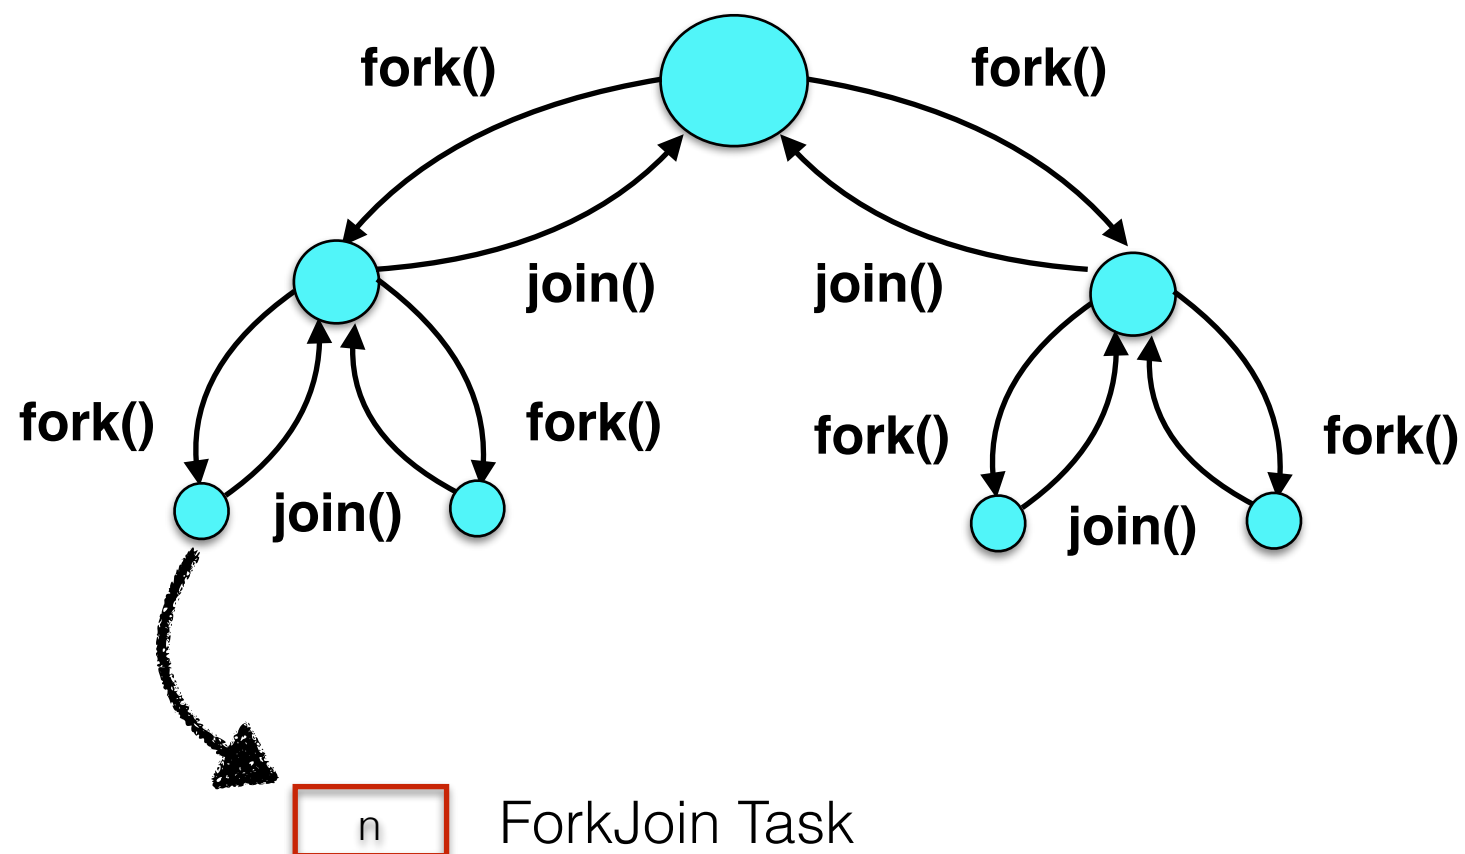
Divider and conquer algorithm

Why
ForkJoin?




Divider and conquer algorithm

Why
ForkJoin?




Why ForkJoin?

Work Stealing




ForkJoin Task




ForkJoin Worker

Why ForkJoin?

Work Stealing



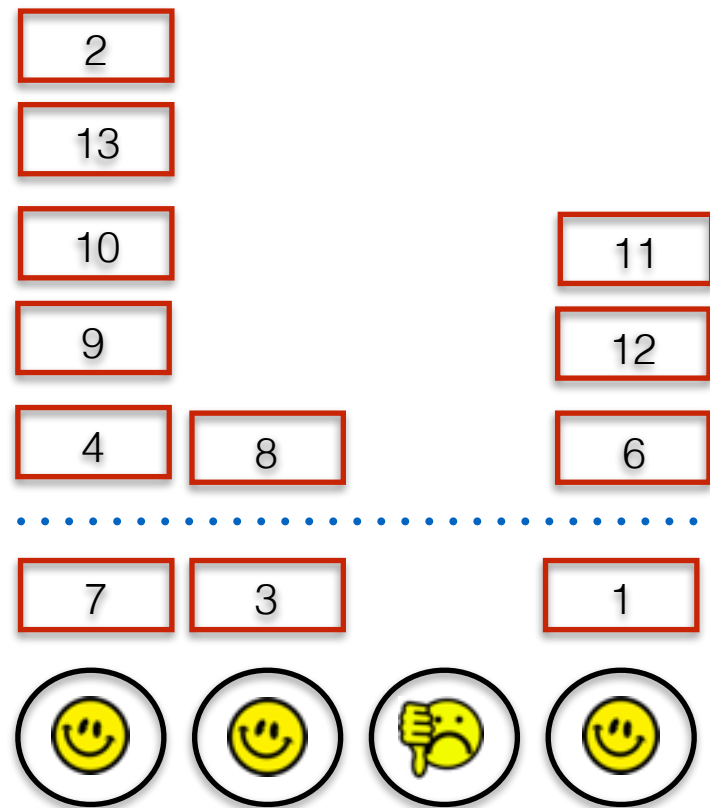
ForkJoin Task



ForkJoin Worker

Why ForkJoin?

Work Stealing



n

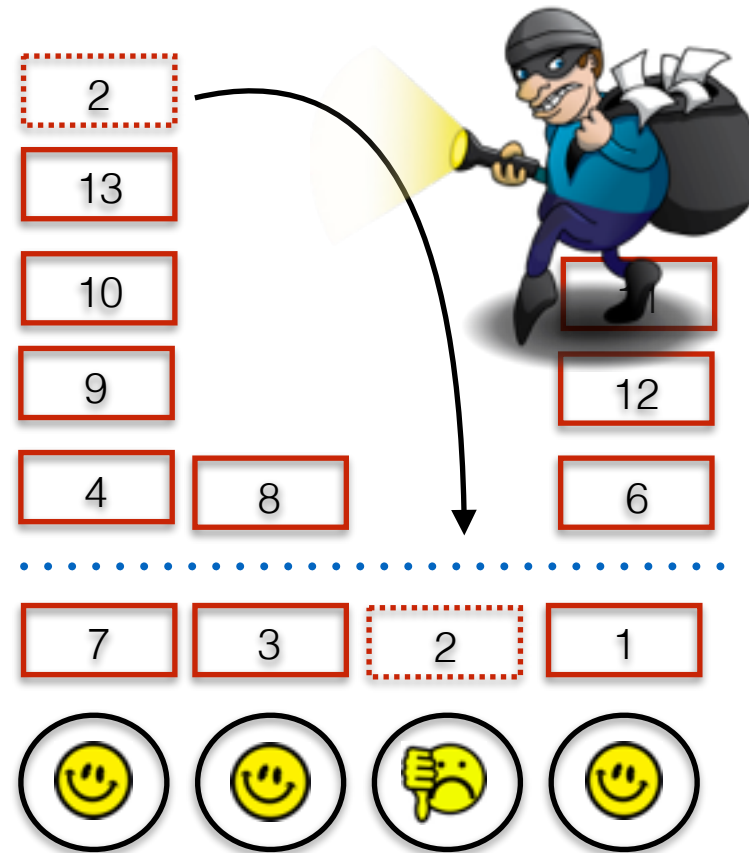
ForkJoin Task

n


ForkJoin Worker

Why ForkJoin?

Work Stealing




ForkJoin Task




ForkJoin Worker

Work Stealing

Why
ForkJoin?





ForkJoin Task



ForkJoin Worker

Bedrock for higher-level Java concurrent libraries

Why ForkJoin?



**System
Programmers**

**Application
Programmers**



**Application
Programmers**

ForkJoin Applications



What are the ForkJoin applications?

1. Search for `java.util.concurrent.ForkJoinPool`
2. Investigate if ForkJoin is indeed used
3. Filter out class assignments and pet projects
4. Try to build and run the code

ForkJoin Applications



What are the ForkJoin applications?

1. Search for `java.util.concurrent.ForkJoinPool`
2. Investigate if ForkJoin is indeed used
3. Filter out class assignments and pet projects
4. Try to build and run the code

30 projects selected

e.g.,

~ 380KLoC

ecco	ejisto	mandelbrot	knn
jacer	conflate	cq4j	mywiki
exhibitor	warp	lowlatency	...

System Programmers

ForkJoin Applications

What about that higher level libs?



1. 330K lines of code
2. 21k commits
3. 470 source code contributors
4. Written (mostly) in Scala and Java
5. Well-known and well-used

This paper



A depth-oriented study and restructuring of the akka message passing algorithm

A breadth-oriented study of 30 real-world ForkJoin open-source projects

A refactoring tool (the first aimed at improving energy consumption of parallel systems)

Understanding Parallelism Bottlenecks



Tree: f7da13b878 FastenTest / src / core / forkjoin / ForkJoinPerformer.java Find file Copy path

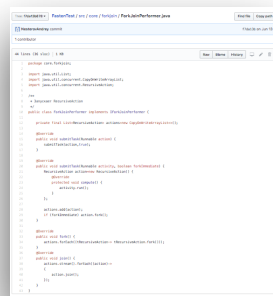
NesterovAndrey commit f7da13b on Jun 13 1 contributor

44 lines (36 sloc) | 1 KB Raw Blame History

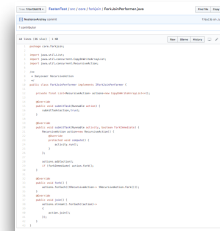


```
1 package core.forkjoin;
2
3 import java.util.List;
4 import java.util.concurrent.CopyOnWriteArrayList;
5 import java.util.concurrent.RecursiveAction;
6
7 /**
8  * Заныкає RecursiveAction
9  */
10 public class ForkJoinPerformer implements IForkJoinPerformer {
11
12     private final List<RecursiveAction> actions=new CopyOnWriteArrayList<>();
13
14     @Override
15     public void submitTask(Runnable action) {
16         submitTask(action,true);
17     }
18
19     @Override
20     public void submitTask(Runnable activity, boolean forkImmediate) {
21         RecursiveAction action=new RecursiveAction() {
22             @Override
23             protected void compute() {
24                 activity.run();
25             }
26         };
27
28         actions.add(action);
29         if (forkImmediate) action.fork();
30     }
31
32     @Override
33     public void fork() {
34         actions.forEach((tRecursiveAction-> tRecursiveAction.fork()));
35     }
36
37     @Override
38     public void join() {
39         actions.stream().forEach((action)->
40         {
41             action.join();
42         });
43     }
```




Understanding Parallelism Bottlenecks



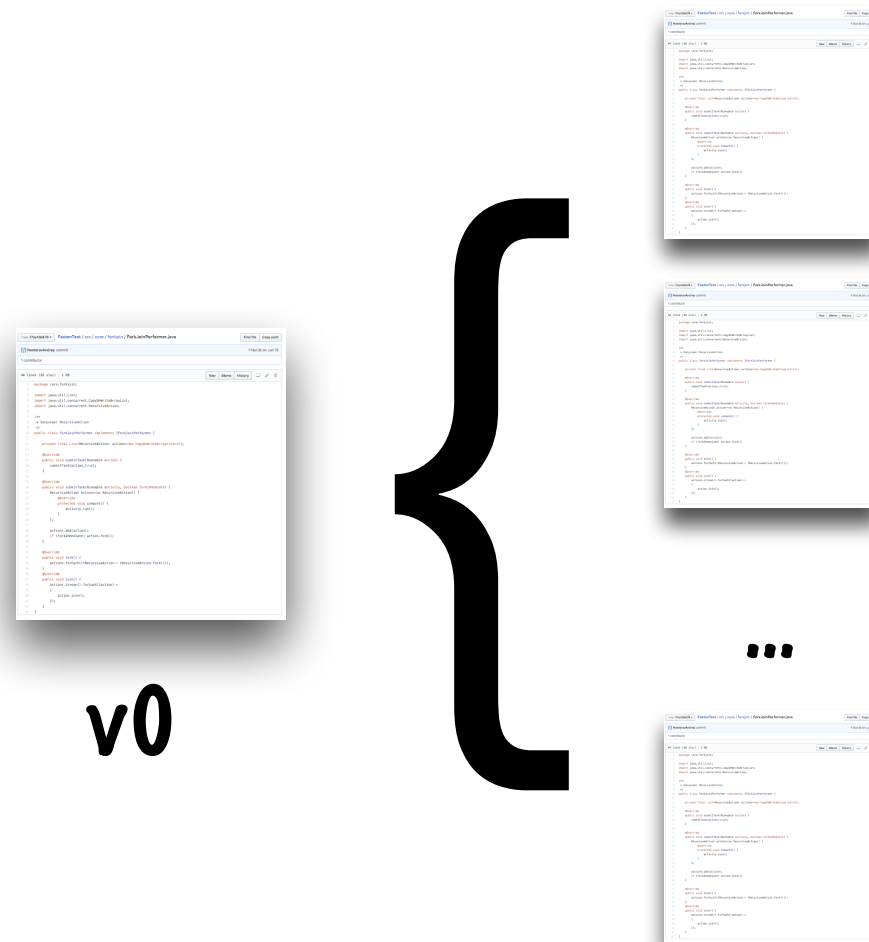
v0



...



Understanding Parallelism Bottlenecks



For each version, we measured execution time
and energy consumption



Intel CPU: A 2×8-core (32-cores w/ hyper-threading),
running Debian, 2.60GHz, with 64GB of memory, JDK
version 1.7.0 71, build 14.

JRapl: Software-based energy measurement

Overcoming Parallelism Bottlenecks

Bottleneck #1: Centralized pooling



actors process their own messages

there is no side effect

Overcoming Parallelism Bottlenecks

Bottleneck #1: Centralized pooling




actors process their own messages

there is no side effect

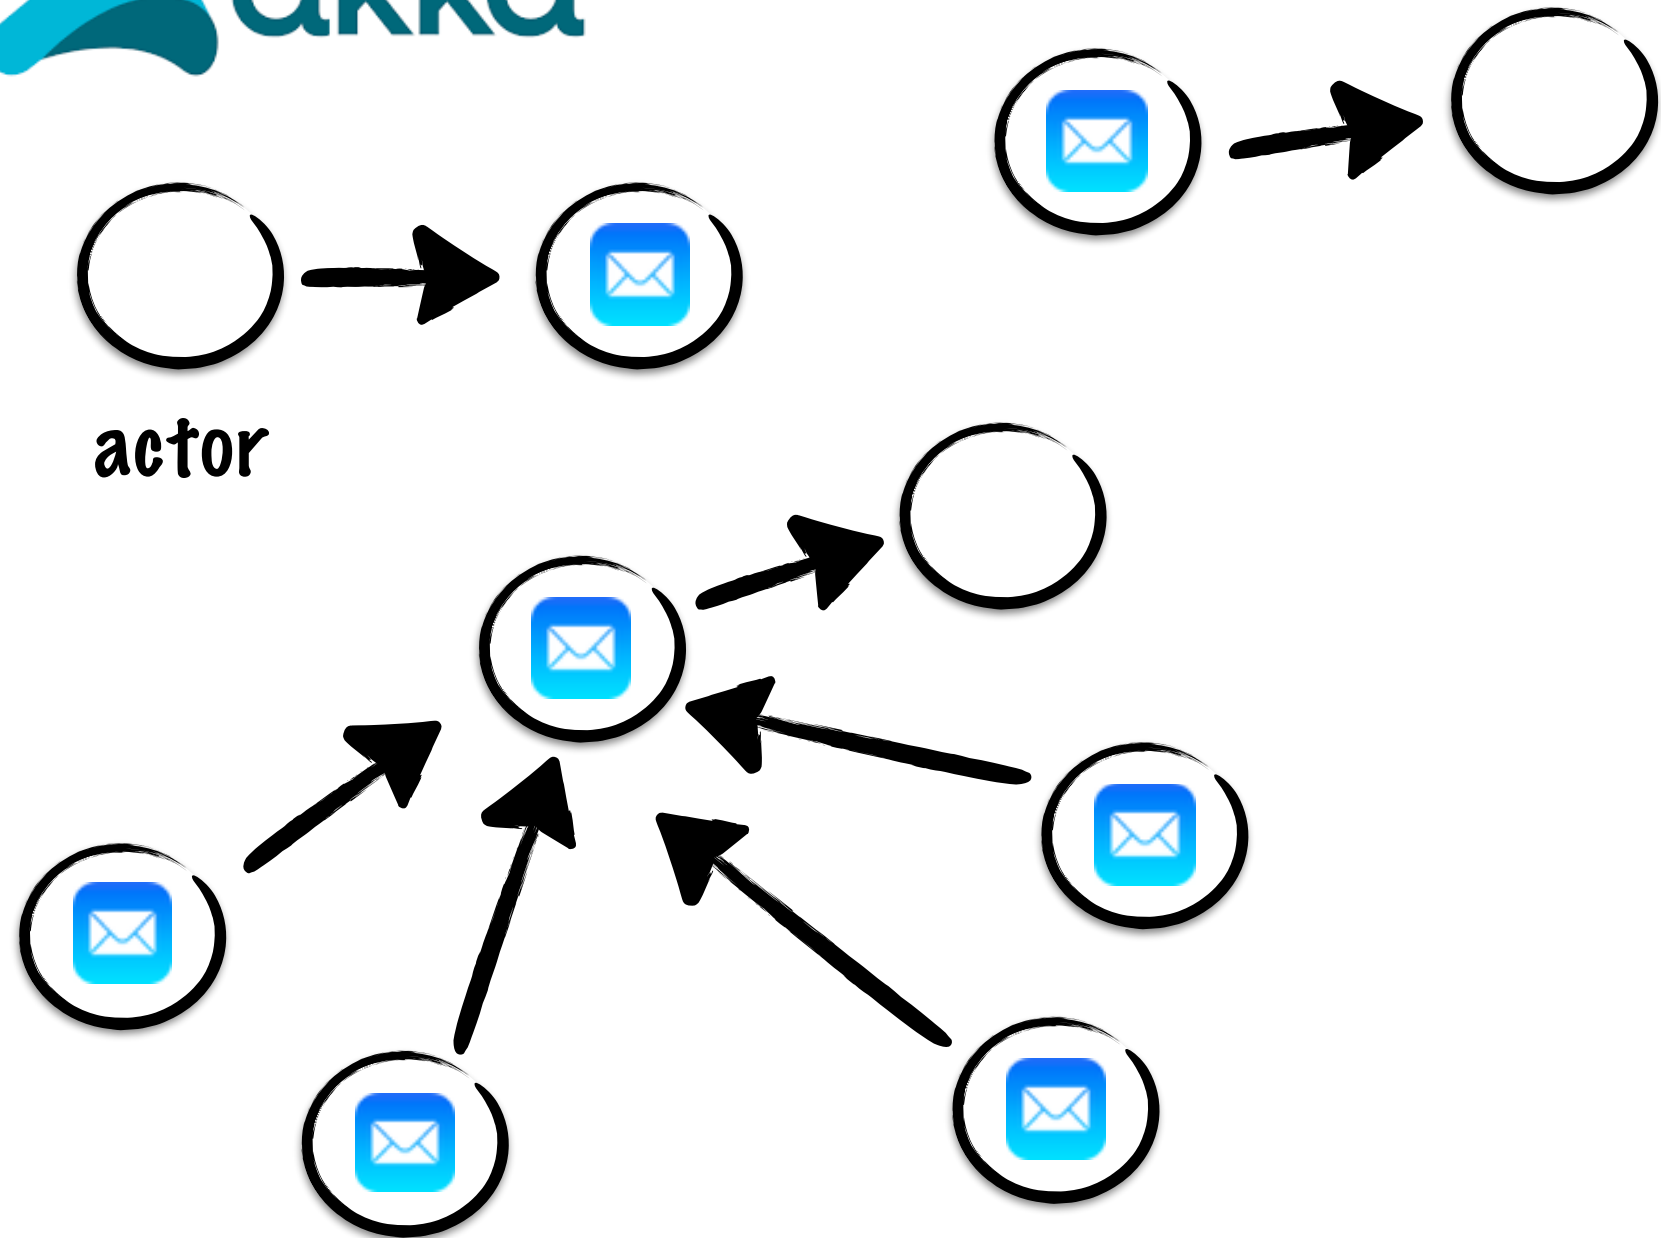
actors exchange, but do not share
the same message

Overcoming Parallelism Bottlenecks

Bottleneck #1: Centralized pooling




Bottleneck #1: Centralized pooling




Overcoming
Parallelism
Bottlenecks

Overcoming Parallelism Bottlenecks

Bottleneck #1: Centralized pooling



actor



actor




Overcoming Parallelism Bottlenecks

Bottleneck #1: Centralized pooling




Overcoming Parallelism Bottlenecks

Bottleneck #1: Centralized pooling




Overcoming Parallelism Bottlenecks

Work Stealing



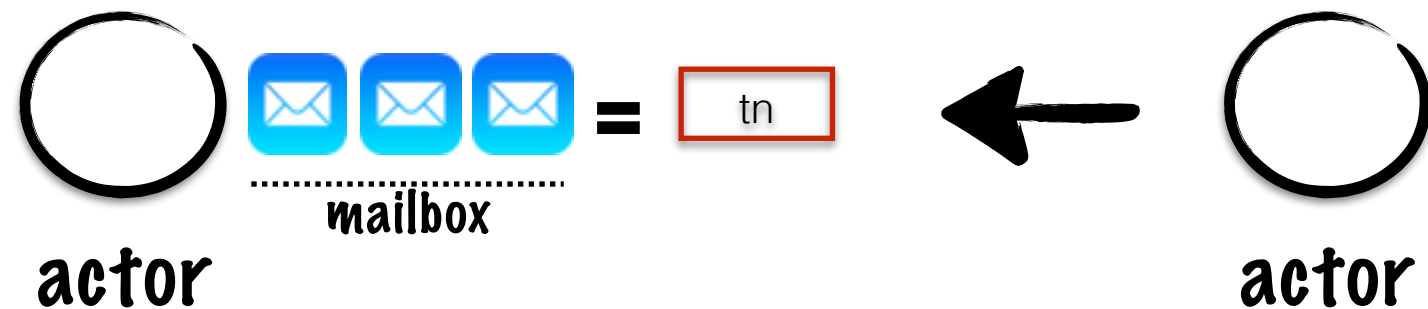
ForkJoin Task



ForkJoin Worker



Overcoming Parallelism Bottlenecks

Bottleneck #1: Centralized pooling





Overcoming Parallelism Bottlenecks

Bottleneck #1: Centralized pooling





Overcoming Parallelism Bottlenecks

Bottleneck #1: Centralized pooling




Overcoming Parallelism Bottlenecks

Bottleneck #1: Centralized pooling




Overcoming Parallelism Bottlenecks

Work Stealing




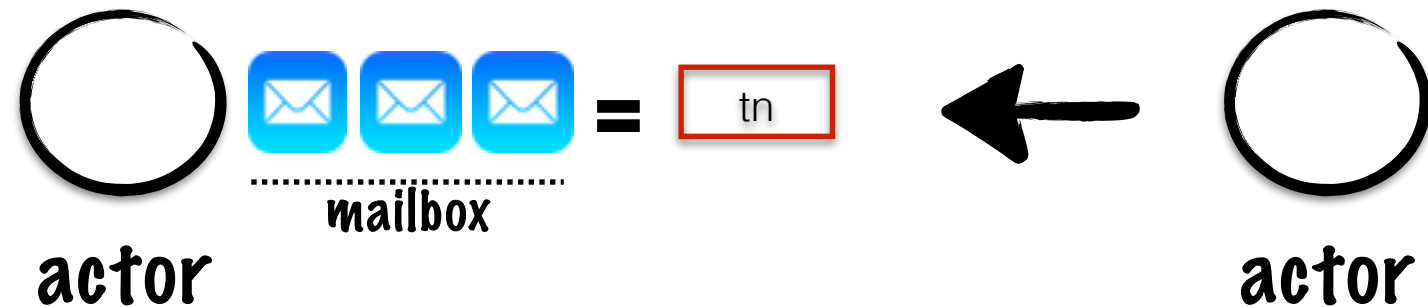
ForkJoin Task



ForkJoin Worker

Overcoming Parallelism Bottlenecks

Bottleneck #1: Centralized pooling



Overcoming Parallelism Bottlenecks

Bottleneck #1: Centralized pooling



benchmark	Runtime (ms)				
	original	σ	custom	σ	speedup
max-throughput	1,861.8	433.4	1,833.0	417.8	1.0×
single-ping	11,657.7	643.8	8,979.6	1,815.5	1.3×
ping-throughput	2,314.9	183.0	701.5	84.7	3.3×
single-producer	4,008.3	1,273.8	4,960.3	2,002.8	0.8×
multi-producer	7,120.3	1,143.9	8,219.5	2,327.0	0.8×
middle-man	3,757.3	195.27	1744.1	195.4	2.1×
mediator	4,633.3	241.13	724.3	123.6	6.4×

Bottleneck #1: Centralized pooling



Overcoming Parallelism Bottlenecks

benchmark	Runtime (ms)				
	original	σ	custom	σ	speedup
max-throughput	1,861.8	433.4	1,833.0	417.8	1.0×
single-ping	11,657.7	643.8	8,979.6	1,815.5	1.3×
ping-throughput	2,314.9	183.0	701.5	84.7	3.3×
single-producer	4,008.3	1,273.8	4,960.3	2,002.8	0.8×
multi-producer	7,120.3	1,143.9	8,219.5	2,327.0	0.8×
middle-man	3,757.3	195.27	1,744.1	195.4	2.1×
mediator	4,633.3	241.13	724.3	123.6	6.4×

Bottleneck #1: Centralized pooling



Overcoming Parallelism Bottlenecks


benchmark	Runtime (ms)				
	original	σ	custom	σ	speedup
max-throughput	1,861.8	433.4	1,833.0	417.8	1.0×
single-ping	11,657.7	643.8	8,979.6	1,815.5	3.3×
ping-throughput	2,314.9	183.0	701.5	84.7	
single-producer	4,008.3	1,273.8	4,960.3	2,002.8	
multi-producer	7,120.3	1,143.9	8,219.5	2,327.0	0.8×
middle-man	3,757.3	195.27	1,744.1	195.4	6.4×
mediator	4,633.3	241.13	724.3	123.6	

Bottleneck #1: Centralized pooling



Overcoming Parallelism Bottlenecks

benchmark	Runtime (ms)				speedup
	original	σ	custom	σ	
max-throughput	1,861.8	433.4	1,833.0	417.8	1.0x
single-ping	11,657.7	643.8	8,979.6	1,815.5	3.3x
ping-throughput	2,314.9	183.0	701.5	84.7	
single-producer	4,008.3	1,273.8	4,960.3	2,002.8	
multi-producer	7,120.3	1,143.9	8,219.5	2,327.0	0.8x
middle-man	3,757.3	195.27	1,744.1	195.4	6.4x
mediator	4,633.3	241.13	724.3	123.6	



Overcoming Parallelism Bottlenecks

Bottleneck #2: Copy on Fork

t1 =

a	b	c	d	e	f	g	h
---	---	---	---	---	---	---	---

t2 = first half

t3 = second half

Overcoming Parallelism Bottlenecks


Bottleneck #2: Copy on Fork

t1 =

a	b	c	d	e	f	g	h
---	---	---	---	---	---	---	---

t2 = first half

t3 = second half



Overcoming Parallelism Bottlenecks


Bottleneck #2: Copy on Fork

t1 =

a	b	c	d	e	f	g	h
---	---	---	---	---	---	---	---

t2 = first half

t3 = second half



Overcoming Parallelism Bottlenecks


Bottleneck #2: Copy on Fork

t1 =

a	b	c	d	e	f	g	h
---	---	---	---	---	---	---	---

t2 = first half

t3 = second half




Bottleneck #2: Copy on Fork

t1 =


a	b	c	d	e	f	g	h
---	---	---	---	---	---	---	---

t2 = first half

t3 = second half



Up to 20% of
energy savings!



t1

t2

0 3

t3

4 7



Overcoming Parallelism Bottlenecks

Bottleneck #3: Copy on Join

t1 =

a	b	c	d
---	---	---	---

t2 =

e	f	g	h
---	---	---	---

t3 = t1 + t2

Overcoming Parallelism Bottlenecks

Bottleneck #3: Copy on Join


t1 =

a	b	c	d
---	---	---	---

t2 =

e	f	g	h
---	---	---	---

t3 = t1 + t2



Overcoming Parallelism Bottlenecks

Bottleneck #3: Copy on Join


t1 =

a	b	c	d
---	---	---	---

t2 =

e	f	g	h
---	---	---	---

t3 = t1 + t2



Bottleneck #3: Copy on Join


t1 =

a	b	c	d
---	---	---	---


t2 =

e	f	g	h
---	---	---	---

t3 = t1 + t2



Up to 12% of energy savings!



Overcoming
Parallelism
Bottlenecks

Overcoming Parallelism Bottlenecks

Bottleneck #4: Scattered Data

t1 = **a****b****a****b****a****b****a** ...

a = memory copies for a subtask

b = forks the subtask


Overcoming Parallelism Bottlenecks

Bottleneck #4: Scattered Data

t1 = **a****b****a****b****a****b** ...

a = memory copies for a subtask

b = forks the subtask




Overcoming Parallelism Bottlenecks

Bottleneck #4: Scattered Data

t1 = **a****b****a****b****a****b** ...

a = memory copies for a subtask

b = forks the subtask




Overcoming Parallelism Bottlenecks

Bottleneck #4: Scattered Data

t1 = **a****b****a****b****a****b** ...

a = memory copies for a subtask

b = forks the subtask



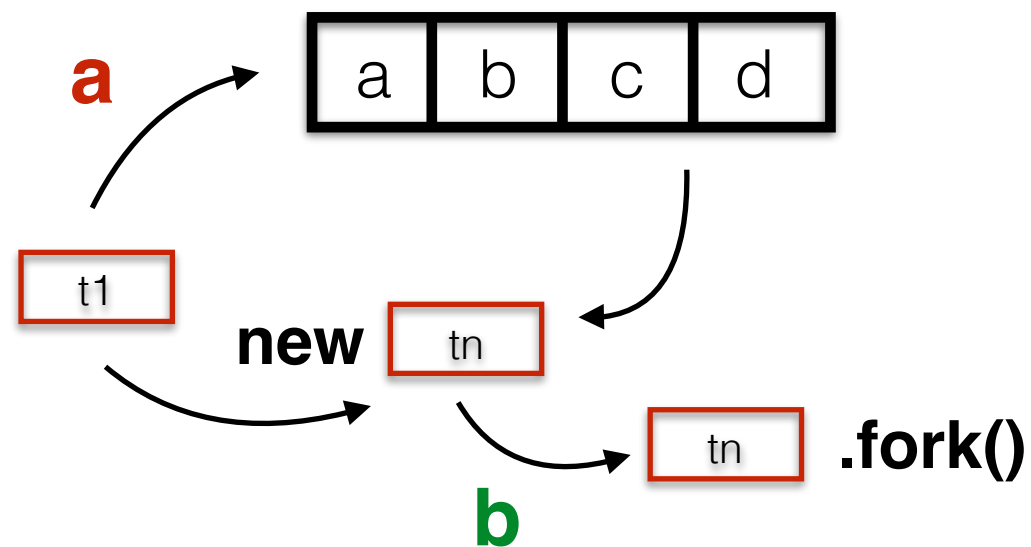
Overcoming Parallelism Bottlenecks

Bottleneck #4: Scattered Data

t1 = **a****b****a****b****a****b** ...

a = memory copies for a subtask

b = forks the subtask



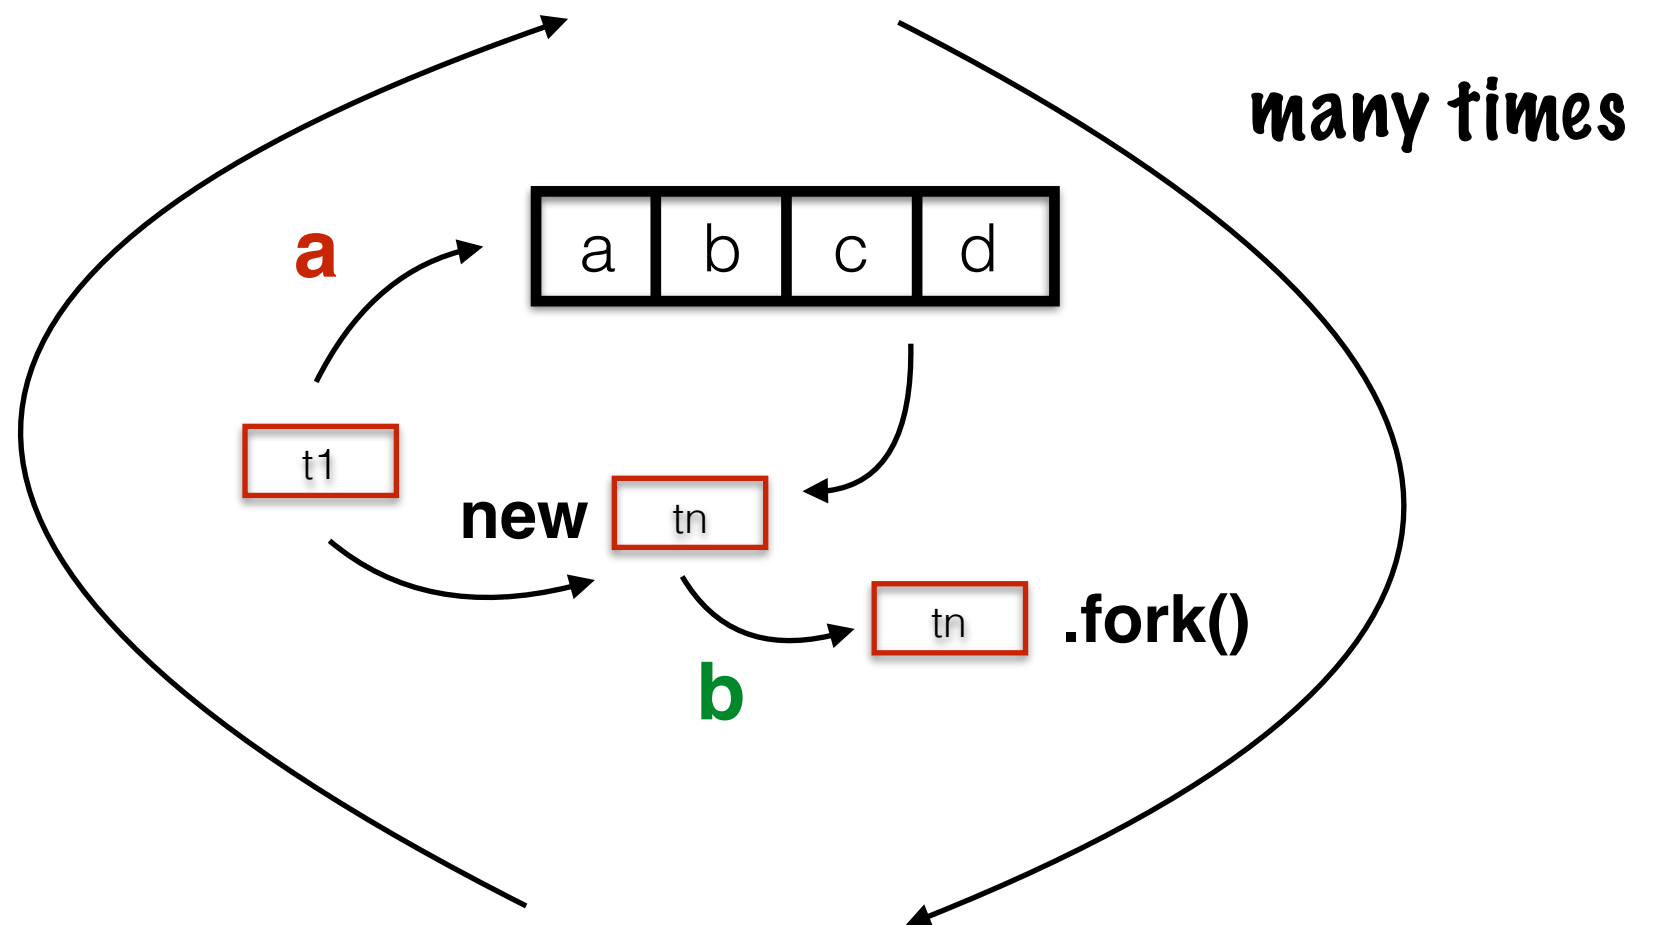
Overcoming Parallelism Bottlenecks

Bottleneck #4: Scattered Data

t1 = **a****b****a****b****a****b** ...

a = memory copies for a subtask

b = forks the subtask



Overcoming Parallelism Bottlenecks

Bottleneck #4: Scattered Data

t1 = **a****b****a****b****a****b****a** ...

a = memory copies for a subtask

b = forks the subtask

t1 = **a****a****a****b****b****b****b**...

Overcoming Parallelism Bottlenecks


Bottleneck #4: Scattered Data

t1 = **a****b****a****b****a****b** ...

a = memory copies for a subtask

b = forks the subtask

t1 = **aaa****bbb**...



Overcoming Parallelism Bottlenecks


Bottleneck #4: Scattered Data

t1 = **a****b****a****b****a****b** ...

a = memory copies for a subtask

b = forks the subtask

t1 = **aaa****bbb**...



Overcoming Parallelism Bottlenecks


Bottleneck #4: Scattered Data

t1 = **a****b****a****b****a****b** ...

a = memory copies for a subtask

b = forks the subtask

t1 = **aaaabbbb**...



Overcoming Parallelism Bottlenecks


Bottleneck #4: Scattered Data

t1 = **a****b****a****b****a****b** ...

a = memory copies for a subtask

b = forks the subtask

t1 = **aaa****bbb**...



after creating the objects..

for task in list:

b **tn**.fork()

Overcoming Parallelism Bottlenecks

Bottleneck #4: Scattered Data

t1 = **a****b****a****b****a****b** ...

a = memory copies for a subtask


b = forks the subtask

t1 = **aaa****bbb**...

10% of energy savings

3% less cache misses

5% less context switches



after creating the objects..

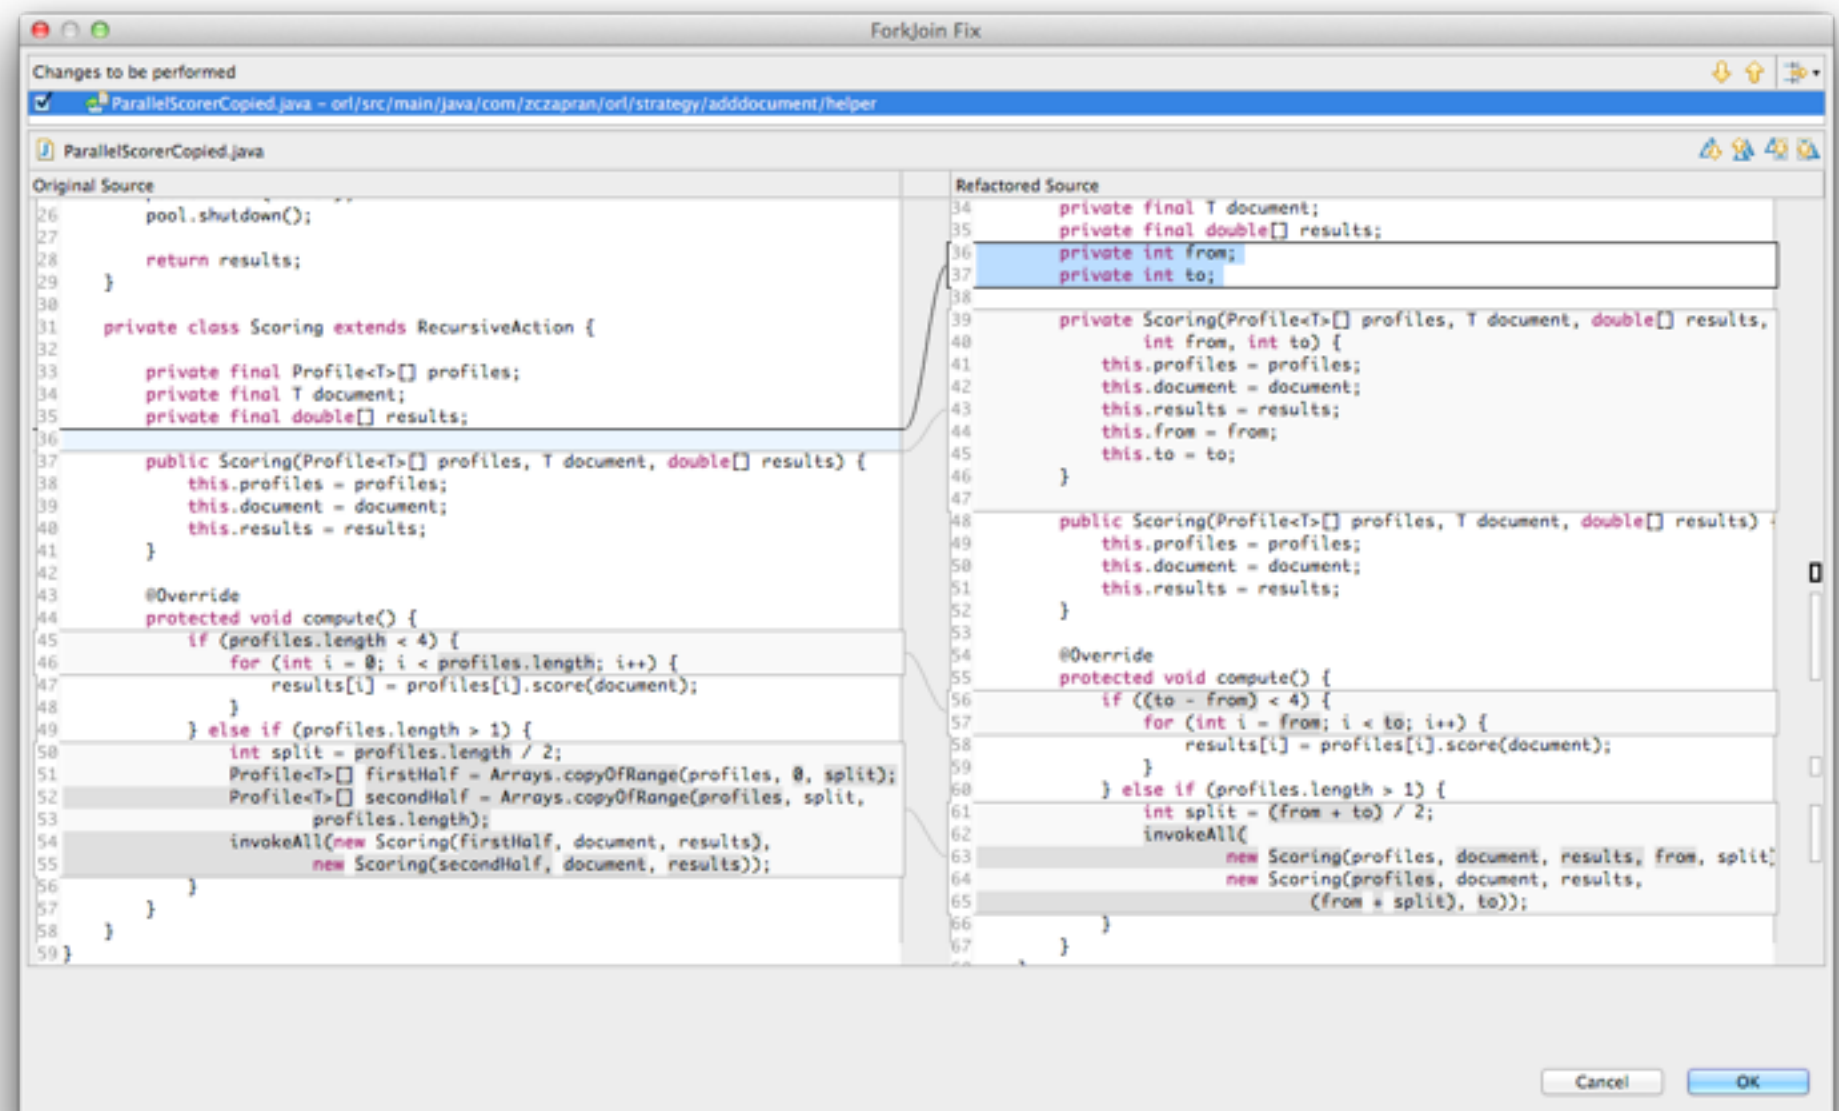
for task in list:

b `tn.fork()`



Automating Bottleneck #2: Copy on Fork

Overcoming Parallelism Bottlenecks



Overcoming Parallelism Bottlenecks

Patching Bottleneck #2: Copy on Fork

Proposed	#1 Improving ForkJoin usage for better performance on Dec 5, 2014 czbabl/itemupdown
Merged	#1 Improving ForkJoin usage for better performance on Dec 5, 2014 toby1984/jAcer
Proposed	#1 Improving ForkJoin usage for better performance on Dec 5, 2014 nikkrichko/Educational
Merged	#1 Improving ForkJoin usage for better performance on Dec 5, 2014 mayukh42/scalatuts
Merged	#1 Improving ForkJoin usage for better performance on Dec 5, 2014 tfreese/knn
Closed	#1 Improving ForkJoin usage for better performance on Dec 5, 2014 TarantulaTechnology/r
Proposed	#1 Improving ForkJoin usage for better performance on Dec 5, 2014 statsbiblioteket/doms-t
Merged	#1 Improving ForkJoin usage for better performance on Dec 5, 2014 2sbsbsb/ForkAndJoin

Proposed	#206 Improving ForkJoin usage for better performance on Dec 5, 2014 Netflix/exhibitor
Proposed	#1 Improving ForkJoin usage for better performance on Dec 5, 2014 robitobi/Solitaire
Merged	#1 Improving ForkJoin usage for better performance on Dec 5, 2014 mariofts/javaOneBR-2012
Proposed	#1 Improving ForkJoin usage for better performance on Dec 5, 2014 cacing/mywiki
Merged	#14 Improving ForkJoin usage for better performance on Dec 5, 2014 ejisto/ejisto
Proposed	#1 Improving ForkJoin usage for better performance on Dec 5, 2014 ctpahhik/cq4j
Merged	#1 Updating ForkJoin usage for better performance and energy on Dec 5, 2014 cjarose/MagicSquares

Overcoming Parallelism Bottlenecks


Patching Bottleneck #2: Copy on Fork

Proposed	#1 Improving ForkJoin usage for better performance on Dec 5, 2015 czbabl/itemupdown
Merged	#1 Improving ForkJoin usage for better performance on Dec 5, 2015 toby1984/jAcer
Proposed	#1 Improving ForkJoin usage for better performance on Dec 5, 2015 nikkrichko/Educational
Merged	#1 Improving ForkJoin usage for better performance on Dec 5, 2015 mayukh42/scalatuts
Merged	#1 Improving ForkJoin usage for better performance on Dec 5, 2015 tfreese/knn
Closed	#1 Improving ForkJoin usage for better performance on Dec 5, 2015 TarantulaTechnology/r
Proposed	#1 Improving ForkJoin usage for better performance on Dec 5, 2015 statsbiblioteket/doms-t
Merged	#1 Improving ForkJoin usage for better performance on Dec 5, 2015 2sbsbsb/ForkAndJoin
Proposed	#206 Improving ForkJoin usage for better performance on Dec 5, 2015 Netflix/exhibitor
Proposed	#1 Improving ForkJoin usage for better performance on Dec 5, 2015 robitobi/Solitaire
Merged	#1 Improving ForkJoin usage for better performance on Dec 5, 2015 mariofts/javaOneBR-2012
Proposed	#1 Improving ForkJoin usage for better performance on Dec 5, 2015 cacing/mywiki
Merged	#14 Improving ForkJoin usage for better performance on Dec 5, 2015 ejisto/ejisto
Proposed	#1 Improving ForkJoin usage for better performance on Dec 5, 2015 ctpahhik/cq4j
Merged	#1 Updating ForkJoin usage for better performance and energy efficiency on Dec 5, 2015 cjlrose/MagicSquares

7/9 of projects that **replied** have **accepted** the PR

Why ForkJoin?

Work Stealing



Understanding Parallelism Bottlenecks



For each version, we measured execution time and energy consumption

Intel CPU: A 2x8-core (32-cores w/ hyper-threading), running Debian, 2.60GHz, with 64GB of memory. JDK version 1.7.0_71, build 14.

JKapl: Software-based energy measurement

@gustavopinto

Overcoming Parallelism Bottlenecks

Bottleneck #4: Scattered Data

t1 = ababababab...

a = memory copies for a subtask


b = forks the subtask

t1 = aaaabbbb...

10% of energy savings

3% less cache misses

5% less context switches



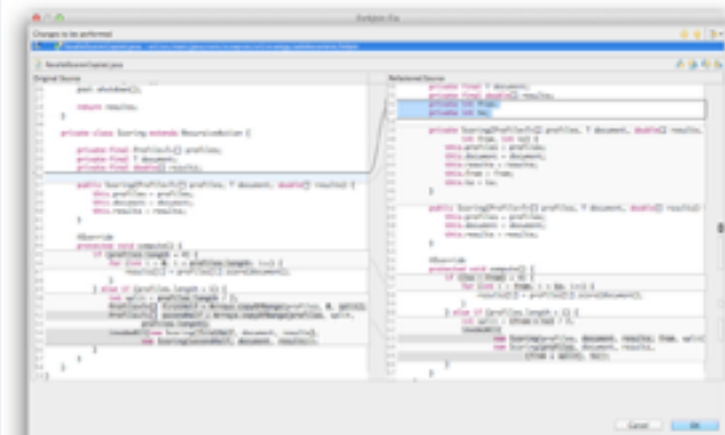
after creating the objects..

for task in list:
tn.fork()



Overcoming Parallelism Bottlenecks


Automating Bottleneck #2: Copy on Fork



@gustavopinto

Why ForkJoin?

Work Stealing



Understanding Parallelism Bottlenecks



For each version, we measured execution time

Questions?



@gustavopinto



gpinto@ufpa.br

Bottleneck #4: Scattered Data

Automating Bottleneck #2: Copy on Fork

Overcoming Parallelism Bottlenecks

t1 = **a**aaaa**b**bbb...

a = memory copies for a subtask


b = forks the subtask

t1 = **aaa****a****bbb****b**...

10% of energy savings

3% less cache misses

5% less context switches

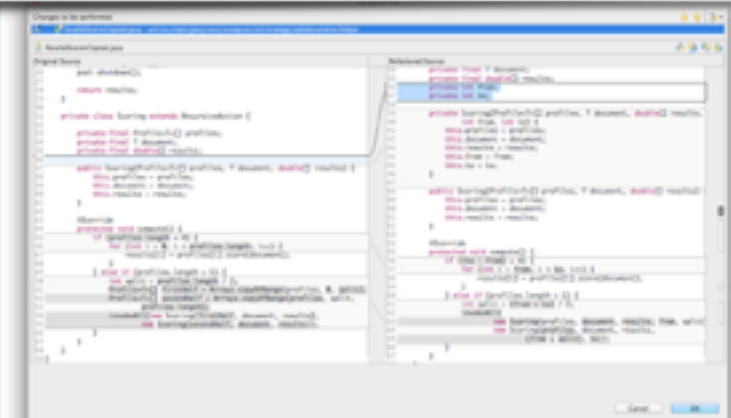


after creating the objects..

for task in list:
 tn.fork()



Overcoming Parallelism Bottlenecks



@gustavopinto

Understanding and Overcoming Parallelism Bottlenecks in ForkJoin Applications



 [@gustavopinto](#)



A. Canino



F. Castor



G. Xu



Y. D. Liu

