

# Mining Energy-Aware Commits



Irineu Moura   Gustavo Pinto   Felipe Ebert   Fernando Castor  
{imlm2, ghlp, fe, castor}@cin.ufpe.br



# The Problem

- Energy efficiency is becoming a key design consideration.
- A large body of work in hardware/architecture, operating systems, runtime systems.



# The Problem



- Energy efficiency is becoming a key design consideration.
- A large body of work in hardware/architecture, operating systems, runtime systems.
- However, little is known about the application level.
- Little is known about what are the **solutions** that programmers actually **use** to mitigate this problem.



4Mi+ Users

19Mi+ Repositories

“GitHub is the largest code host on the planet with over 19.9 mi repositories.”


<https://github.com/features>



19M Repos




19M Repos



Query on  
**GitHub Archive**



19M Repos




Query on  
GitHub Archive

```
SELECT (repository_url + '/commit/' + payload_commit_id) as commit_link, repository_url,
repository_owner, repository_name, payload_commit_id, payload_commit_email,
payload_commit_msg, payload_commit_name, payload_commit_flag, created_at
FROM [githubarchive:github.timeline]
WHERE payload_commit_msg is not null
AND PARSE_UTC_USEC(created_at) <= PARSE_UTC_USEC ('2014-05-15 23:59:59')
AND (lower(payload_commit_msg) like '%power consum%'
OR lower(payload_commit_msg) like '%power efficien%'
OR lower(payload_commit_msg) like '%power sav%'
OR lower(payload_commit_msg) like '%save power%'
OR lower(payload_commit_msg) like '%energy consum%'
OR lower(payload_commit_msg) like '%energy efficien%'
OR lower(payload_commit_msg) like '%energy sav%'
OR lower(payload_commit_msg) like '%save energy%')
AND type = 'PushEvent'
GROUP BY commit_link, payload_commit_msg, repository_url, payload_commit_id,
created_at, repository_owner, repository_name, payload_commit_email,
payload_commit_name, payload_commit_flag
ORDER BY created_at asc
```





19M Repos



Query on  
GitHub Archive


2,189 Commits

```
SELECT (repository_url + '/commit/' + payload_commit_id) as commit_link, repository_url,
       repository_owner, repository_name, payload_commit_id, payload_commit_email,
       payload_commit_msg, payload_commit_name, payload_commit_flag, created_at
FROM [githubarchive:github.timeline]
WHERE payload_commit_msg is not null
AND PARSE_UTC_USEC(created_at) <= PARSE_UTC_USEC('2014-05-15 23:59:59')
AND (lower(payload_commit_msg) like '%power consum%'
     OR lower(payload_commit_msg) like '%power efficien%'
     OR lower(payload_commit_msg) like '%power sav%'
     OR lower(payload_commit_msg) like '%save power%'
     OR lower(payload_commit_msg) like '%energy consum%'
     OR lower(payload_commit_msg) like '%energy efficien%'
     OR lower(payload_commit_msg) like '%energy sav%'
     OR lower(payload_commit_msg) like '%save energy%')
AND type = 'PushEvent'
GROUP BY commit_link, payload_commit_msg, repository_url, payload_commit_id,
         created_at, repository_owner, repository_name, payload_commit_email,
         payload_commit_name, payload_commit_flag
ORDER BY created_at asc
```





19M Repos



Query on  
**GitHub Archive**



Automatic Filter


1,005 Commits




19M Repos



Automatic Filter



Manual Filter



Query on  
GitHub **Archive**






19M Repos



Automatic Filter



Manual Filter

### Commits

dh7h3

md8ja

j287h

dij873

dj827h

os837

82uan

28a08

2ja82

d0hk0

j29yd

a7jf9

aio92

hnna2






19M Repos



Automatic Filter



Manual Filter

### Commits

dh7h3

md8ja

j287h

dij873

dj827h

os837

82uan

28a08

2ja82

d0hk0

j29yd

a7jf9

aio92

hnna2






19M Repos



Automatic Filter



Manual Filter

### Commits

dh7h3

md8ja

j287h

dij873

dj827h

os837

82uan

28a08

2ja82

d0hk0

j29yd

a7jf9

aio92

hnna2






19M Repos



Automatic Filter



Manual Filter

### Commits

dh7h3

md8ja

j287h

dij873

dj827h

os837

82uan

28a08

2ja82

d0hk0

j29yd

a7jf9

aio92

hnna2



All commits  
were, at least,  
double-checked!







19M Repos



Automatic Filter



Manual Filter







Query on  
GitHub Archive



371  
Energy-Aware  
Commits

# 371 Energy-Aware Commits

- From 03/12/2012 to 05/15/2014
- Performed into 317 OSS Projects
- Performed by 265 OSS Developers

Metric	Mean	Median	Standard Dev.	Histogram
LoC	4,069,000	68,930	5,239,099	
Contributors	202.3	6.5	335.31	
Commits	68,250	475.5	126,104	
Age	4.50	3	3.72	

# Research Questions

- **RQ1.** What are the **solutions** that developers use to save energy in practice?
- **RQ2.** What software **quality attributes** may be given precedence over energy consumption?
- **RQ3.** How are energy-saving solutions distributed over the **software stack**?
- **RQ4.** To what extent are software developers **certain** that their commits will save energy?

# RQ1: Solutions

- Frequency and voltage scaling (50 occurrences)
- Use power efficient library/device (45 occurrences)
- Disabling features or devices (42 occurrences)
- Energy bug fix (26 occurrences)
- Low power idling (22 occurrences)
- Timing out (16 occurrences)

# RQ1: Solutions

- Frequency and voltage scaling (50 occurrences)

### Reduce Wifi voltage for power savings

Should be beneficial for a wifi only device

10.1

 Metallice authored on Jun 27, 2012

1 parent [3cee2da](#) commit [34f4738019a9e04c882ffa05bed1b40b81017c51](#)

Showing 1 changed file with 1 addition and 1 deletion.

Unified Split

2  arch/arm/mach-omap2/board-espresso-wifi.c

View

	@@ -230,7 +230,7 @@	static struct regulator_init_data espresso_vmmc5 = {
230	230	
231	231	static struct fixed_voltage_config espresso_vwlan = {
232	232	.supply_name = "vwl1271",
233	-	.microvolts = 2000000, /* 2.0V */
	233	+ .microvolts = 1800000, /* 2.0V */ /*1.8V Metallice*/
234	234	.startup_delay = 70000, /* 70msec */
235	235	.enable_high = 1,
236	236	.enabled_at_boot = 0,

# RQ1: Solutions

- Use power efficient library/device (45 occurrences)

**TEST: WCD9310: Use power efficient workqueue**[Browse files](#)

Convert our WCD9310 codec to use the new power efficient workqueue offered by newer kernels.

master

yseras authored on Feb 17, 2014 1 parent [08a755b](#) commit [7618d681692c6ae1508bd981c7bc837b8f2bc334](#)

Showing **1 changed file** with **4 additions** and **4 deletions**. Unified Split

8 sound/soc/codecs/wcd9310.c [View](#)

			@@ -2912,7 +2912,7 @@ static int tabla_codec_enable_dec(struct snd_soc_dapm_widget *w,
2912	2912		if (tx_hpf_work[decimator - 1].tx_hpf_cut_of_freq !=
2913	2913		CF_MIN_3DB_150HZ) {
2914	2914		
2915		-	schedule_delayed_work(&tx_hpf_work[decimator - 1].dwork,
	2915	+	queue_delayed_work(system_power_efficient_wq, &tx_hpf_work[decimator - 1].dwork,
2916	2916		msecs_to_jiffies(300));
2917	2917		}



# RQ2: Quality Attributes

- Correctness (7 occurrences)
- Responsiveness (6 occurrences)
- Performance (3 occurrences)
- No actual power saving (3 occurrences)
- Miscellaneous (3 occurrences)


# RQ2: Quality Attributes

- Correctness (7 occurrences)

## semi working - sensors read out, but bad radio

[Browse files](#)

Quashed compile bug and redid some documentation, but radio does not transmit proper values – it probably can't handle the array in the typedef struct. will have to fix that. the sensors scan fine and write correctly to the array (varified by serial output) – also, the power saving delay has the ability to corrupt serial transmissions. changing these to regular delays fixes this. – getting there!

 master (#4)

 **Teslafly** authored on Apr 19, 2014

1 parent [b5bc8fa](#) commit [659afda362c0e0a23dc36d6f56fa0aab83812c91](#)

# RQ2: Quality Attributes

- Responsiveness (6 occurrences)

## cpufreq: boost the sampling rate on touch event


[Browse files](#)

For better Ux responsiveness on demand sampling rate needs to be 20ms. But, a 20ms sampling rate increases power consumption. So, boost the sampling rate to 20ms on every touch event for 2.5 ms and later reset to default rate. Also, change sampling down factor proportional to the sampling rate.

Change-Id: I111b1cf3b8ed133347149afc34d329d0384ecfcb

Signed-off-by: Narayanan Gopalakrishnan <nargop@codeaurora.org>

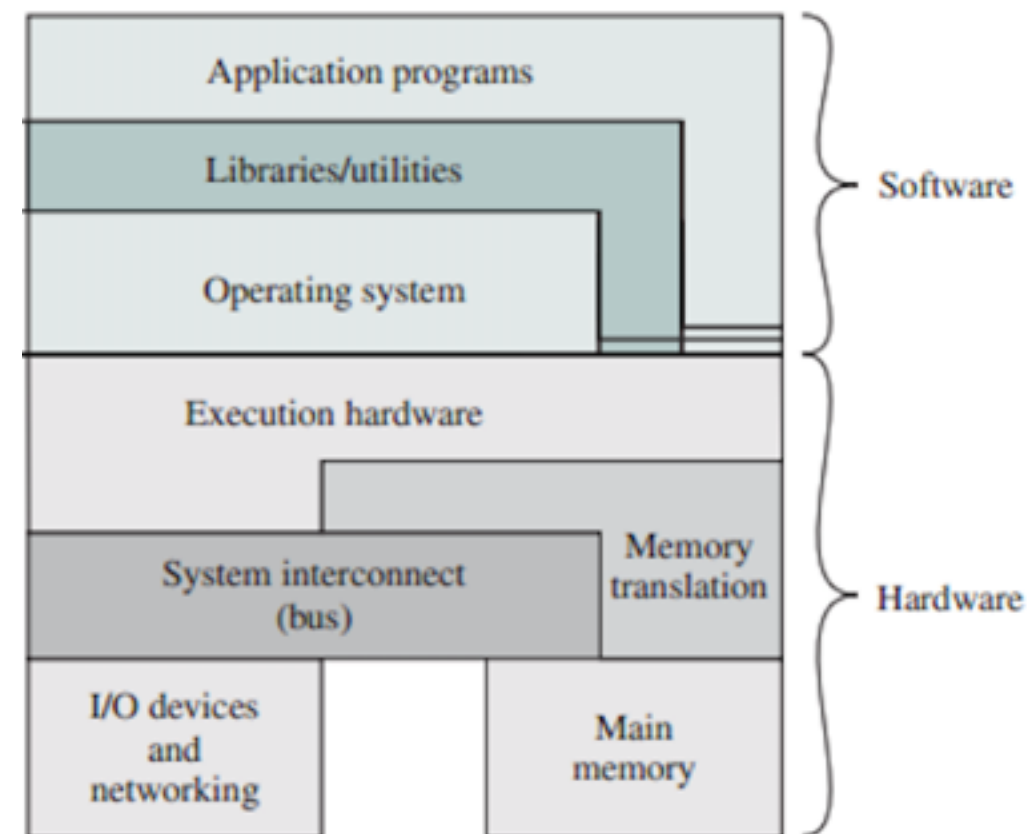

cm-10.1\_3.0 + jellybean + jellybean-stable + mr1-staging

 Narayanan Gopalakrishnan authored on Aug 9, 2012

1 parent [b56043a](#) commit [694447e65dd4af2af7193021ab67c4cad91ff412](#)

→ [Intervigillum](#) committed on Oct 5, 2012


# RQ3: Software Stack



# RQ3: Software Stack

88 Commits ←

Application includes  
embedded applications,  
desktop application, and  
mobile applications.

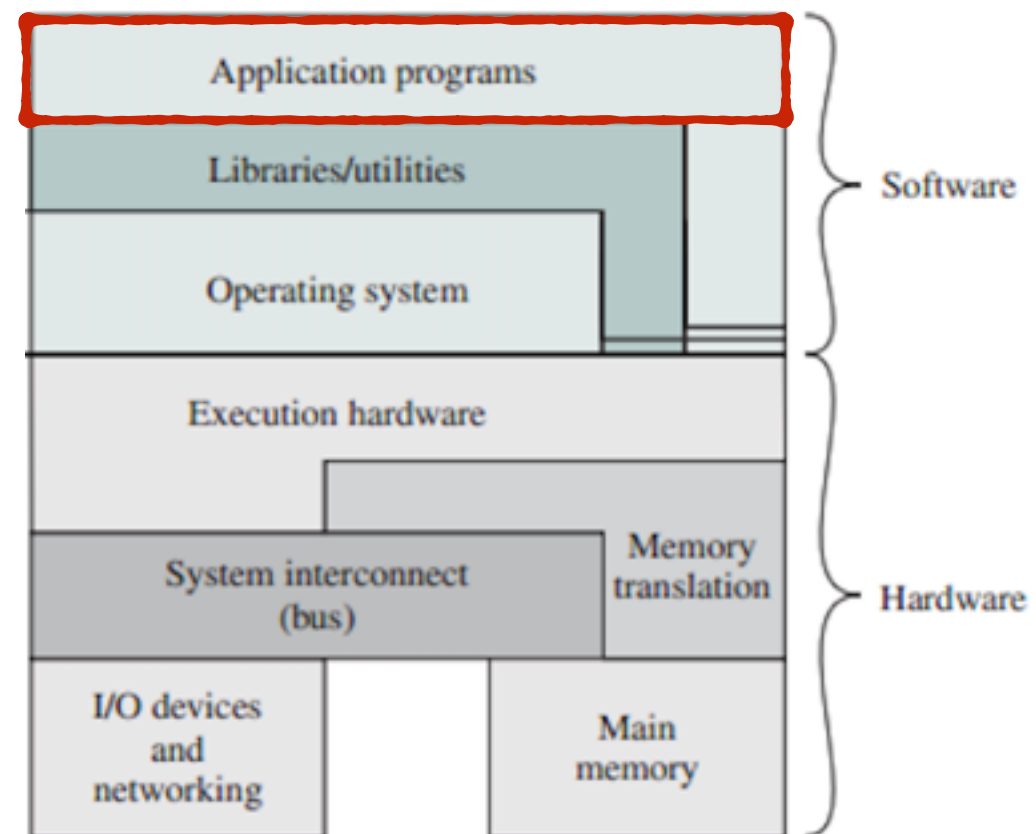


# RQ3: Software Stack

88 Commits ←

42 Embedded

Application includes  
embedded applications,  
desktop application, and  
mobile applications.





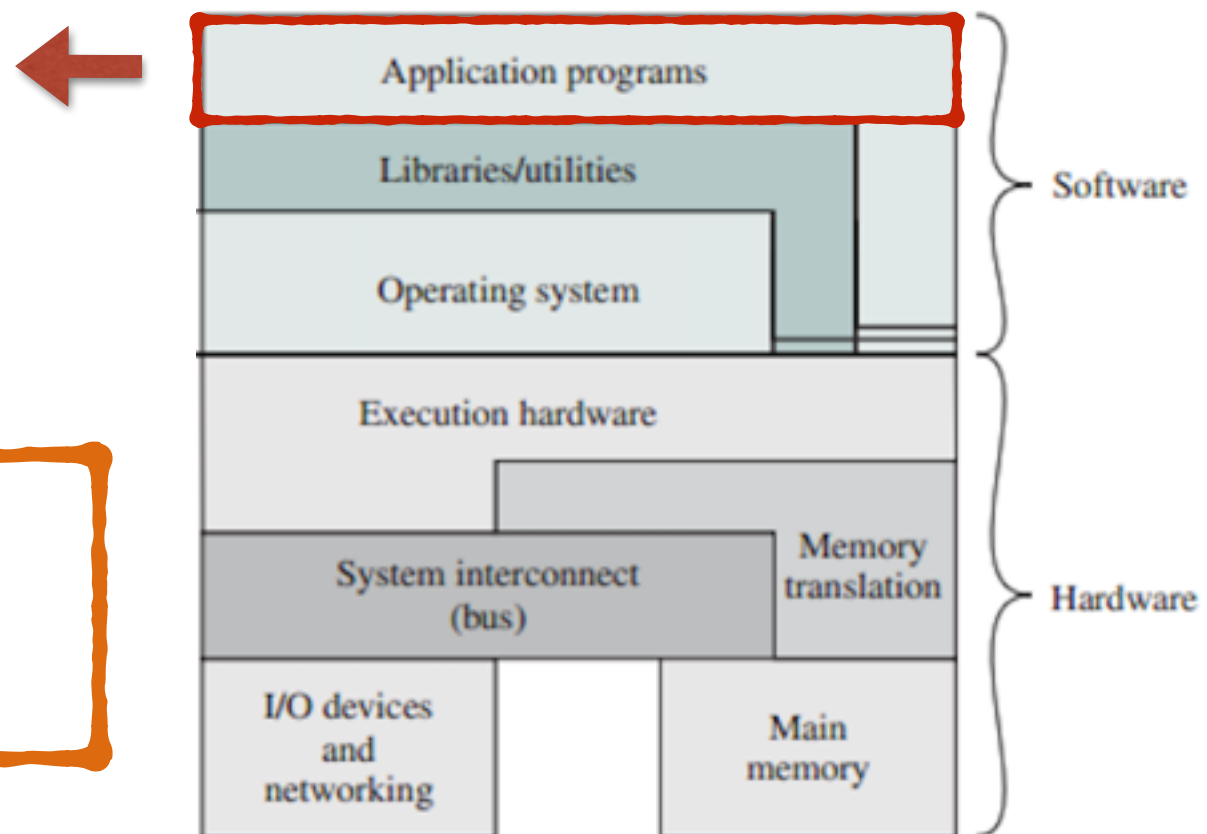
# RQ3: Software Stack

88 Commits

42 Embedded


21 Arduino

Application includes  
embedded applications,  
desktop application, and  
mobile applications.

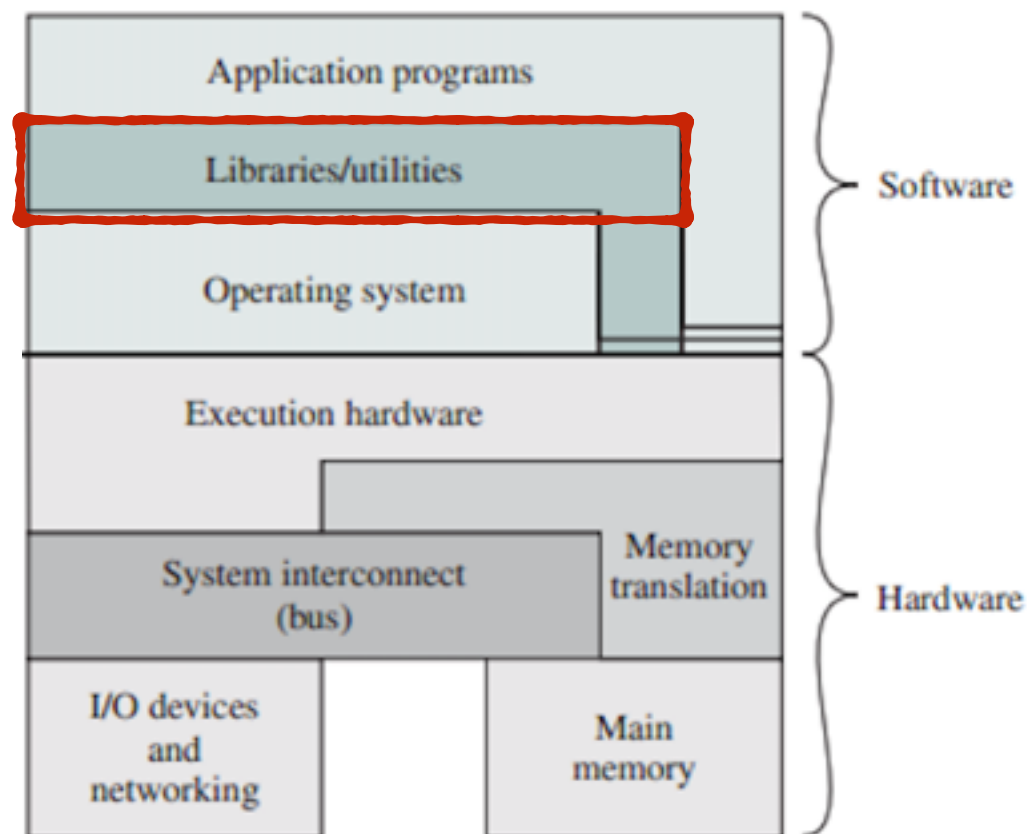


# RQ3: Software Stack

50 Commits



Libraries/Utilities include scripts and embedded libraries.

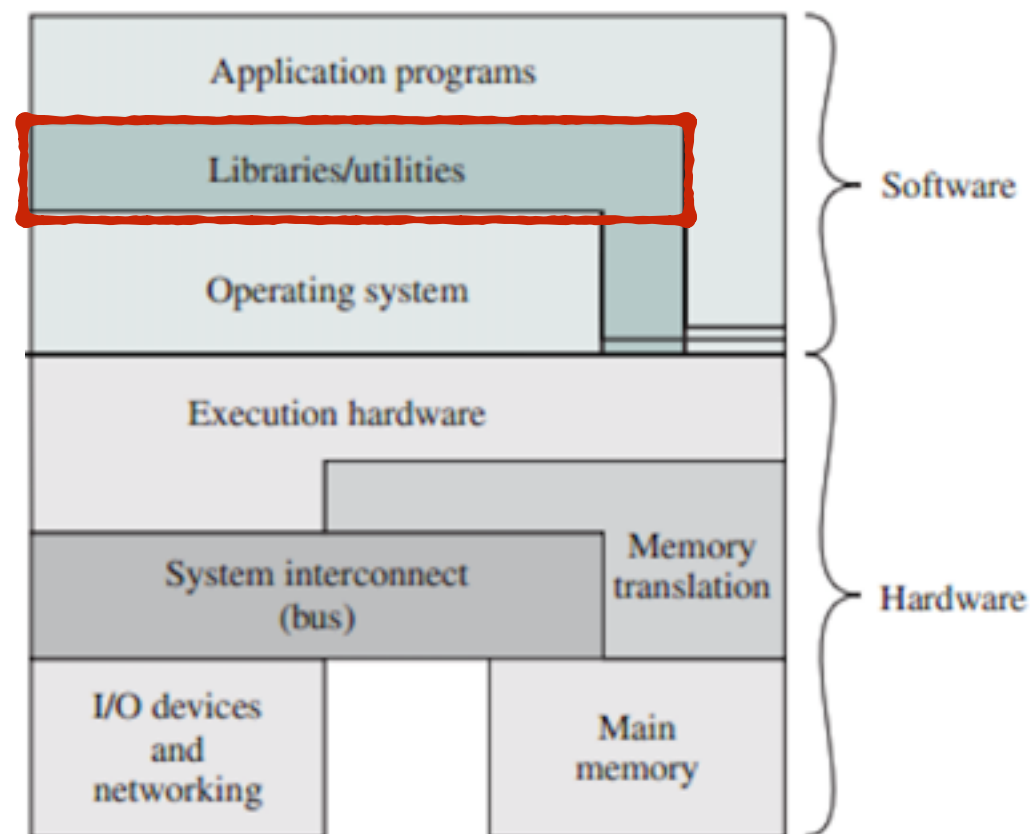


# RQ3: Software Stack

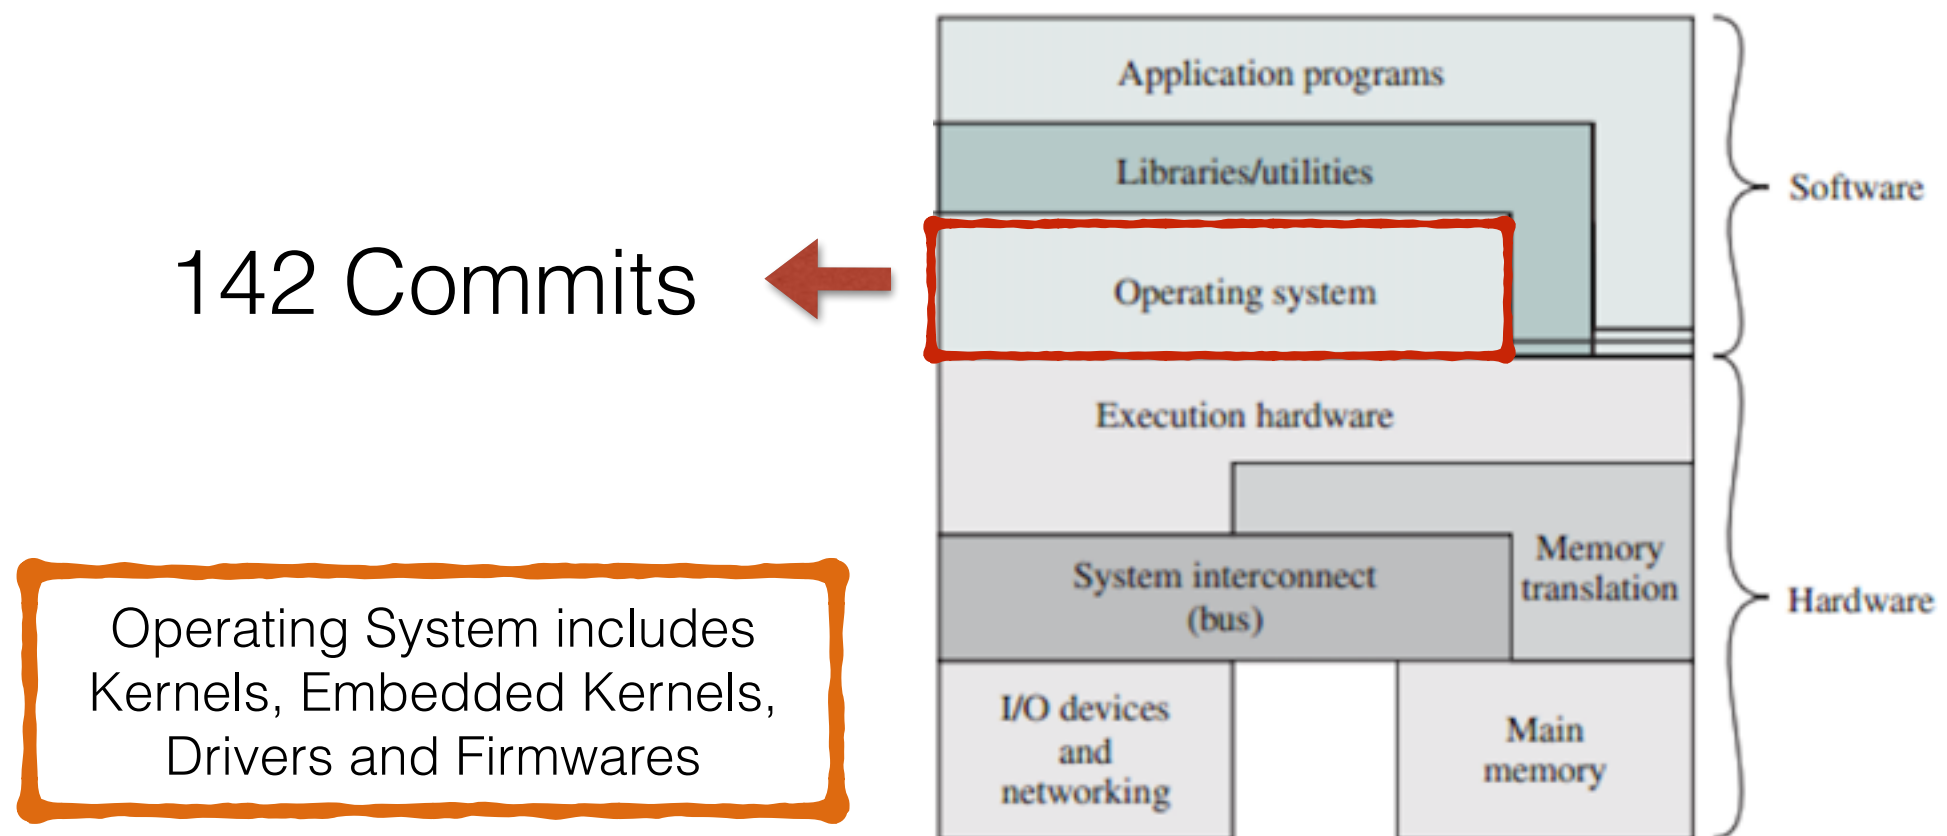
50 Commits

41 scripts

Libraries/Utilities include scripts  
and embedded libraries.



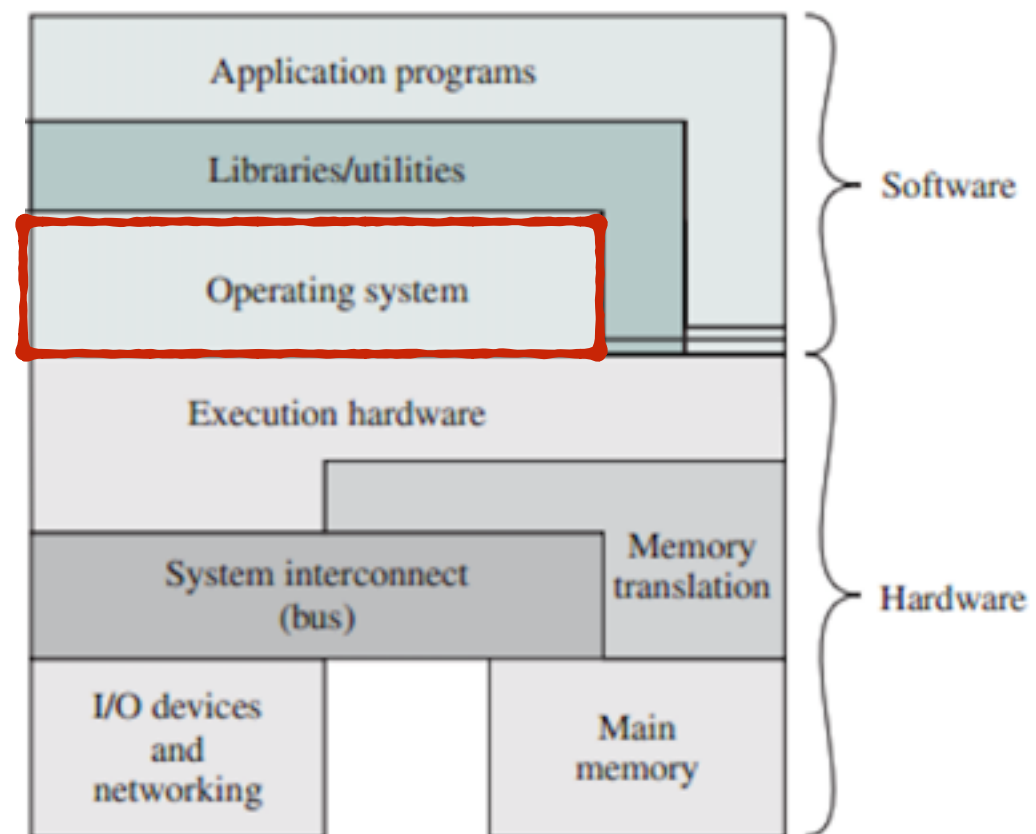
# RQ3: Software Stack



# RQ3: Software Stack

142 Commits  
69 Kernels

Operating System includes  
Kernels, Embedded Kernels,  
Drivers and Firmwares



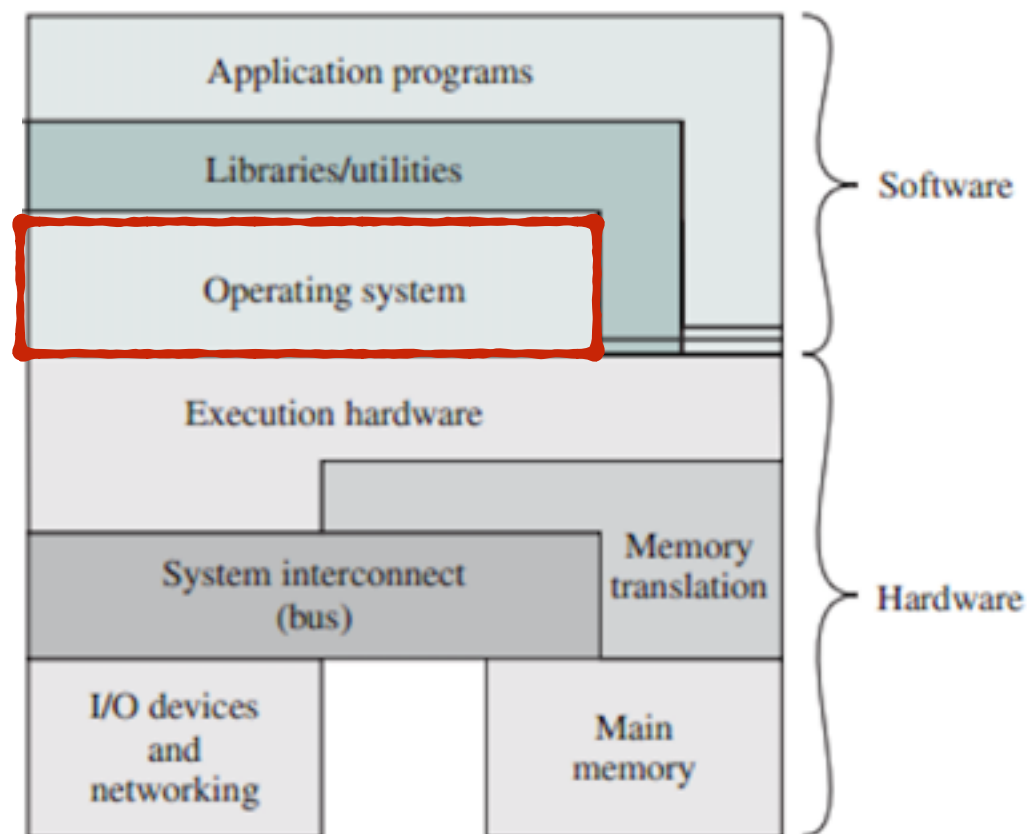
# RQ3: Software Stack

142 Commits

69 Kernels

54 Drivers

Operating System includes  
Kernels, Embedded Kernels,  
Drivers and Firmwares





# RQ4: Certain

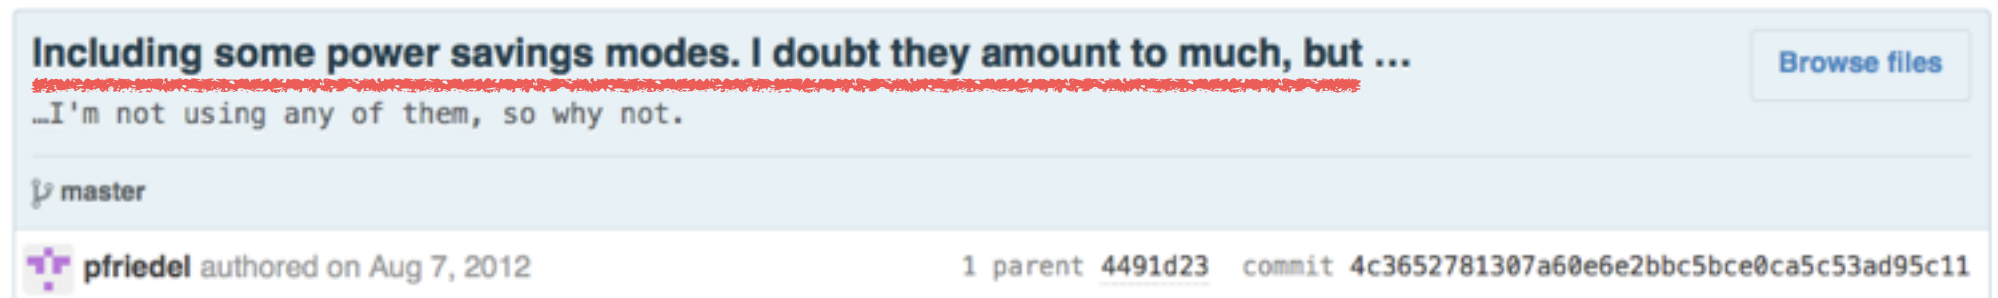
“Hesitating” words

- seem
- might
- doubt
- could
- hope
- attempt
- supposed
- guess
- likely

# RQ4: Certain

“Hesitating” words      18 hesitations!

- seem
- might
- doubt
- could
- hope
- attempt
- supposed
- guess
- likely



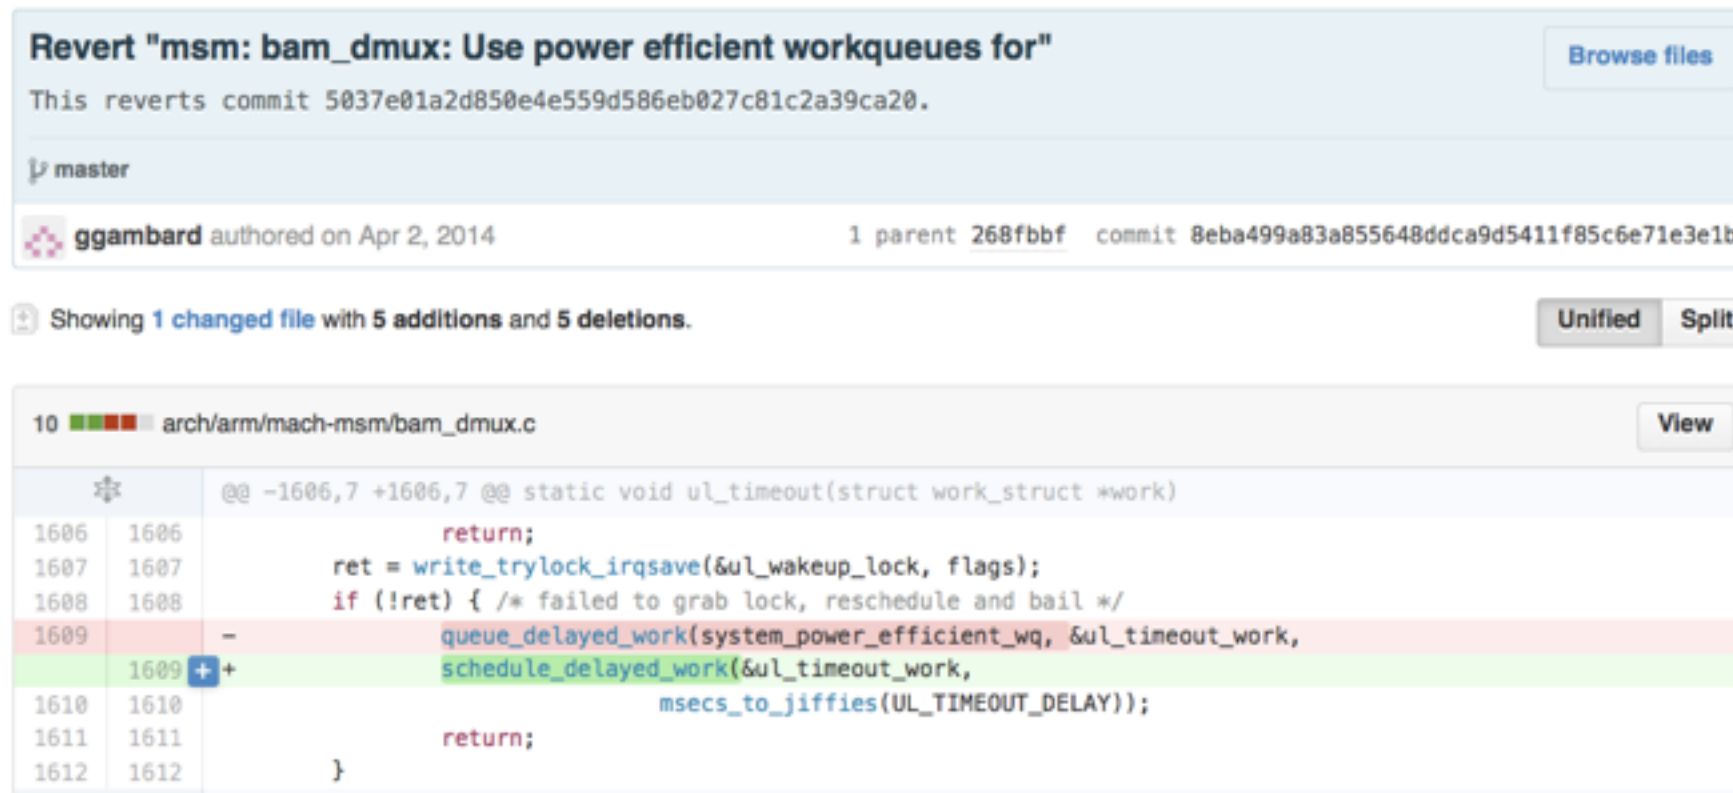
# RQ4: Certain

18 Reverted Commits

# RQ4: Certain

18 Reverted Commits

8/18 reverts the power efficient work queue!



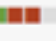
**Revert "msm: bam\_dmux: Use power efficient workqueues for"** [Browse files](#)

This reverts commit 5037e01a2d850e4e559d586eb027c81c2a39ca20.

master

ggambard authored on Apr 2, 2014 1 parent 268fbbf commit 8eba499a83a855648ddca9d5411f85c6e71e3e1b

Showing 1 changed file with 5 additions and 5 deletions. [Unified](#) [Split](#)

10  arch/arm/mach-msm/bam\_dmux.c [View](#)

```
@@ -1606,7 +1606,7 @@ static void ul_timeout(struct work_struct *work)
1606 1606         return;
1607 1607         ret = write_trylock_irqsave(&ul_wakeup_lock, flags);
1608 1608         if (!ret) { /* failed to grab lock, reschedule and bail */
1609 -         queue_delayed_work(system_power_efficient_wq, &ul_timeout_work,
1609 +         schedule_delayed_work(&ul_timeout_work,
1610                               msecs_to_jiffies(UL_TIMEOUT_DELAY));
1611 1611         return;
1612 1612     }
```

# RQ4: Certain

18 Reverted Commits

8/18 reverts the power efficient work queue!

There is no  
silver bullet!

**Revert "msm: bam\_dmux: Use power efficient workqueues for"** [Browse files](#)

This reverts commit 5037e01a2d850e4e559d586eb027c81c2a39ca20.

master

ggambard authored on Apr 2, 2014 1 parent 268fbbf commit 8eba499a83a855648ddca9d5411f85c6e71e3e1b

Showing 1 changed file with 5 additions and 5 deletions. [Unified](#) [Split](#)

10 arch/arm/mach-msm/bam\_dmux.c [View](#)

	@@ -1606,7 +1606,7 @@	static void ul_timeout(struct work_struct *work)
1606	1606	return;
1607	1607	ret = write_trylock_irqsave(&ul_wakeup_lock, flags);
1608	1608	if (!ret) { /* failed to grab lock, reschedule and bail */
1609	-	queue_delayed_work(system_power_efficient_wq, &ul_timeout_work,
	1609 +	schedule_delayed_work(&ul_timeout_work,
1610	1610	msecs_to_jiffies(UL_TIMEOUT_DELAY));
1611	1611	return;
1612	1612	}




Fred Brooks

# Conclusions

## RQ1: Solutions

- Frequency and voltage scaling (50 occurrences)



16

## RQ2: Quality Attributes


- Correctness (7 occurrences)
- Responsiveness (6 occurrences)
- Performance (3 occurrences)
- No actual power saving (3 occurrences)
- Miscellaneous (3 occurrences)

27

## RQ3: Software Stack

50 Commits  
41 scripts

Libraries/Utilities include scripts  
and embedded libraries.



36

## RQ4: Certain

"Hesitating" words


- seem
- might
- could
- hope
- attempt
- guess
- ...



40

[bit.ly/energy-aware-mining](http://bit.ly/energy-aware-mining)

# Mining Energy-Aware Commits



Irineu Moura   Gustavo Pinto   Felipe Ebert   Fernando Castor  
{imlm2, ghlp, fe, castor}@cin.ufpe.br

